

Snowflake Performance Optimization and Cost Management Strategy

1. Introduction

Snowflake is a cloud-native data platform that provides scalable storage and compute resources. While Snowflake simplifies data warehousing operations, inefficient queries and poor warehouse management can lead to increased costs and degraded performance.

This report demonstrates:

- Query profiling using Snowflake Query Profiler.
- Query optimization techniques.
- Performance improvements through clustering and filtering.
- Warehouse management best practices.
- Cost monitoring and optimization strategies.

2. Sample Dataset Overview

Sales Transactions Table

The sample dataset consists of a large sales transaction table containing historical transaction records.

Table Structure

Column Name	Data Type
TRANSACTION_ID	NUMBER
CUSTOMER_ID	NUMBER
PRODUCT_ID	NUMBER
REGION	STRING
ORDER_DATE	DATE
QUANTITY	NUMBER
SALES_AMOUNT	NUMBER(12,2)

Dataset Characteristics

Metric	Value
Total Rows	100 Million
Table Size	500 GB
Historical Data	5 Years
Regions	12


3. Query Profiling Analysis

Original Query

```
SELECT
  c.CUSTOMER_NAME, p.PRODUCT_NAME, SUM(s.SALES_AMOUNT) as TOTAL_SALES
FROM SALES_TRANSACTIONS s
JOIN CUSTOMERS c
  ON s.CUSTOMER_ID = c.CUSTOMER_ID
JOIN PRODUCTS p
  ON s.PRODUCT_ID = p.PRODUCT_ID
WHERE YEAR(s.ORDER_DATE) = 2024
GROUP BY
  c.CUSTOMER_NAME,
  p.PRODUCT_NAME;
```

Query Profiler Results

Metric	Value
Execution Time	48 Seconds
Bytes Scanned	420 GB
Partitions Scanned	98%
Warehouse Size	Medium
Credits Used	0.40



Performance Issues Identified

1. Function-based filtering on ORDER_DATE prevents partition pruning.
2. Excessive data scanning due to poor filtering.
3. Lack of clustering strategy.
4. Unnecessary join operations increase execution overhead.
5. High warehouse resource consumption.

4. Query Optimization Process

Optimization 1: Date Range Filtering

Before

WHERE YEAR(ORDER_DATE) = 2024

After

```
WHERE ORDER_DATE >= '2024-01-01'  
AND ORDER_DATE < '2025-01-01'
```

Benefits

- Enables micro-partition pruning.
- Reduces data scanning.
- Improves query execution speed.

Optimization 2: Implement Clustering Keys

Clustering Key Script

```
ALTER TABLE SALES_TRANSACTIONS  
CLUSTER BY (ORDER_DATE, REGION);
```

Validation

```
SELECT SYSTEM$CLUSTERING_INFORMATION(  
'SALES_TRANSACTIONS'  
);
```

Benefits

- Improved partition elimination.
- Faster range-based queries.
- Reduced I/O operations.

Optimization 3: Eliminate Redundant Joins

When product information is not required in the output, the PRODUCTS table join can be removed.


Optimized Query

```
SELECT  
  c.CUSTOMER_NAME,  
  SUM(s.SALES_AMOUNT) AS TOTAL_SALES  
FROM SALES_TRANSACTIONS s  
JOIN CUSTOMERS c  
  ON s.CUSTOMER_ID = c.CUSTOMER_ID  
WHERE ORDER_DATE >= '2024-01-01'  
AND ORDER_DATE < '2025-01-01'
```

GROUP BY c.CUSTOMER_NAME;

Benefits

- Reduced join processing.
- Lower memory utilization.
- Faster execution.



5. Performance Results

Before and After Comparison

Metric	Before	After	Improvement
Execution Time	48 sec	11 sec	77%
Bytes Scanned	420 GB	85 GB	80%
Partitions Scanned	98%	18%	82%
Credits Used	0.40	0.10	75%

Analysis

The optimization techniques significantly reduced query execution time and resource consumption. The largest gains were achieved through proper filtering and clustering key implementation.

6. Materialized View Optimization

For frequently accessed aggregated data, a materialized view can improve performance.

Materialized View

```
CREATE MATERIALIZED VIEW MV_MONTHLY_SALES AS
SELECT
    REGION,
    DATE_TRUNC('MONTH', ORDER_DATE) AS SALES_MONTH,
    SUM(SALES_AMOUNT) AS TOTAL_SALES
FROM SALES_TRANSACTIONS
GROUP BY 1,2;
```

Performance Impact

Query Source	Runtime
Base Table	15 sec
Materialized View	2 sec

7. Snowflake Cost Management Strategy

7.1 Warehouse Management Best Practices

Right-Size Warehouses

Use warehouse sizes appropriate for workload requirements.

Workload Type	Recommended Warehouse
Dashboard Queries	X-Small
Reporting	Small
Analytics	Medium
ETL Processing	Large

Enable Auto Suspend

```
ALTER WAREHOUSE ANALYTICS_WH
SET AUTO_SUSPEND = 60;
```

This suspends the warehouse after 60 seconds of inactivity, reducing unnecessary credit consumption.

Enable Auto Resume

```
ALTER WAREHOUSE ANALYTICS_WH
SET AUTO_RESUME = TRUE;
```

This automatically restarts the warehouse when queries are submitted.

7.2 Workload Isolation Strategy

Separate workloads into dedicated warehouses.

Recommended Warehouse Structure:

- ETL_WH
- REPORTING_WH
- BI_WH
- DATA_SCIENCE_WH
- ADHOC_WH

Benefits:

- Prevents resource contention.
- Improves performance consistency.
- Simplifies cost allocation.

7.3 Query Cost Monitoring

Query Cost Monitoring Script

```
SELECT
    QUERY_ID,
    USER_NAME,
    WAREHOUSE_NAME,
    EXECUTION_TIME,
    BYTES_SCANNED,
    CREDITS_USED_CLOUD_SERVICES
FROM SNOWFLAKE.ACCOUNT_USAGE.QUERY_HISTORY
WHERE START_TIME >= DATEADD(DAY,-7,CURRENT_TIMESTAMP());
```

Identify Expensive Queries

```
SELECT
    QUERY_TEXT,
    TOTAL_ELAPSED_TIME,
    BYTES_SCANNED
FROM SNOWFLAKE.ACCOUNT_USAGE.QUERY_HISTORY
ORDER BY BYTES_SCANNED DESC
LIMIT 20;
```

7.4 Resource Monitoring

Create Resource Monitor

```
CREATE RESOURCE MONITOR MONTHLY_BUDGET
WITH CREDIT_QUOTA = 100
FREQUENCY = MONTHLY
TRIGGERS
ON 75 PERCENT DO NOTIFY
ON 90 PERCENT DO NOTIFY
ON 100 PERCENT DO SUSPEND;
```

Assign Resource Monitor

```
ALTER WAREHOUSE ANALYTICS_WH
SET RESOURCE_MONITOR = MONTHLY_BUDGET;
```

Benefits:

- Prevents budget overruns.
- Provides spending alerts.
- Automatically suspends usage when limits are reached.

Benefits:

- Automatically handles concurrent workloads.
- Improves user experience during peak demand.
- Prevents query queuing.

Results (just now)		
Table Chart Pivot		
	WAREHOUSE_NAME	TOTAL_CREDITS
1	COMPUTE_WH	37.325547221
2	TEMP_WAREHOUSE	0.621202499
3	DEMO_WH	0.188042778
4	MY_WAREHOUSE	0.135385833
5	SETUP_WH	0.114646945
6	ANALYTICS_WH	0.098958612
7	TEMP	0.073513056
8	CLOUD_SERVICES_ONLY	0.000260555

8. Cost Calculation Example

Before Optimization

Warehouse Size: Medium

Credits per Hour: 4

Query Runtime: 48 Seconds

Credits Per Query = $(48 \div 3600) \times 4$

Credits Per Query = 0.053

Monthly Usage (10,000 Executions):

$0.053 \times 10,000 = 530$ Credits

After Optimization

Query Runtime: 11 Seconds

Credits Per Query = $(11 \div 3600) \times 4$, Credits Per Query = 0.012

Monthly Usage (10,000 Executions):

$0.012 \times 10,000 = 120$ Credits

Monthly Savings

530 Credits – 120 Credits = 410 Credits Saved

Percentage Reduction: $(410 \div 530) \times 100 = 77\%$

9. Cost Monitoring Dashboard Query

```
SELECT
  DATE_TRUNC('DAY', START_TIME) AS DAY,
  SUM(CREDITS_USED_CLOUD_SERVICES) AS DAILY_CREDITS
FROM SNOWFLAKE.ACCOUNT_USAGE.QUERY_HISTORY
GROUP BY 1
ORDER BY 1;
```

10. Conclusion

The performance optimization exercise demonstrates that strategic query tuning, clustering key implementation, and warehouse management can significantly improve Snowflake performance while reducing operational costs.

Key outcomes include:

- 77% reduction in execution time.
- 80% reduction in bytes scanned.
- 75% reduction in credit consumption.
- Improved query scalability.
- Better workload isolation and governance.
- Enhanced cost visibility through monitoring and resource controls.

By implementing these recommendations, organizations can achieve both performance improvements and sustainable Snowflake cost management.